



# **mlflow** Platform for Complete Machine Learning Lifecycle

Jules S. Damji  
[@2twitme](https://twitter.com/2twitme)

San Francisco | April 29, 2020: Part 1 of 4 Series

\$ whoami




Apache Spark Developer & Community Advocate @ Databricks

Developer Advocate @ Hortonworks

Software engineering @ Sun Microsystems, Netscape, @Home, Excite@Home, VeriSign, Scalix, Centrify, LoudCloud/Opware, ProQuest

Program Co-chair Spark + AI Summit

[@2twitme](https://www.linkedin.com/in/dmatrix)





## VISION

Accelerate innovation by unifying data science, engineering and business to solve data problems

---

## SOLUTION

Unified Data Analytics Platform

---

## WHO WE ARE

- Original creators of  **APACHE Spark**™
  -  **DELTA LAKE**
  -  **mlflow**™
  -  **Koalas**
- 2000+ global companies use our platform across big data & machine learning lifecycle

# Outline – Part 1

- Overview of ML development challenges
- Concepts and Motivations
- How MLflow tackles these
- MLflow Components
  - MLflow Tracking
    - Build and Track metrics, params, runs
    - User MLflow UI to compare runs
- Q & A

# Machine Learning Development is Complex

# Traditional Software

**Goal:** Meet a functional specification

Quality depends only on code

Typically pick one software stack

# Machine Learning

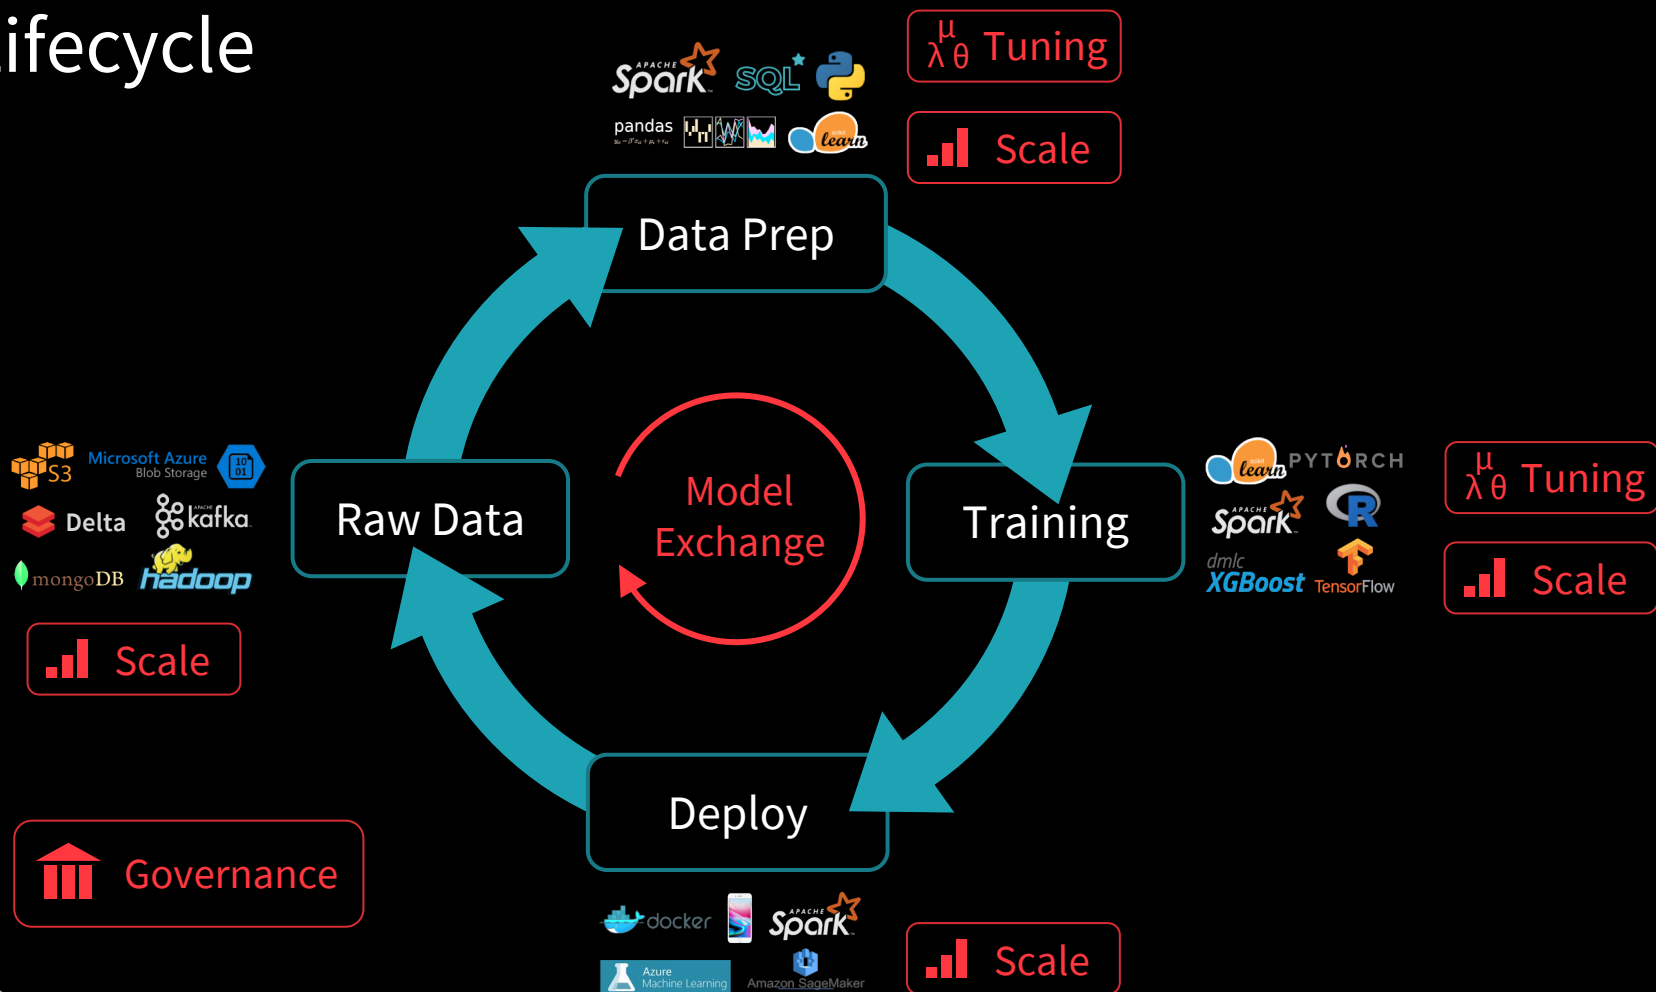
**Goal:** Optimize a metric (e.g., accuracy)  
Constantly experiment to improve it

Quality depends on input data  
and tuning parameters



Compare + combine many libraries,  
models & algorithms for the same task

# ML Lifecycle



# Custom ML Platforms

## Some Big Data Companies

- + **Standardize the data prep / training / deploy loop:**  
if you work with the platform, you get these!
- **Limited to a few algorithms or frameworks**
- **Tied to one company's infrastructure**
- **Out of luck if you left the company....**

**Can we provide similar benefits in an **open** manner?**



# Introducing mlflow

## Open machine learning platform

- Works with popular ML library & language
- Runs the same way anywhere (e.g., any cloud or locally)
- Designed to be useful for 1 or 1000+ person orgs
- *Simple. Modular. Easy-to-use.*
- *Offers positive developer experience to get started!*

# MLflow Design Philosophy

## “API-first”

- Submit runs, log models, metrics, etc. from popular library & language
- Abstract “model” lambda function that MLflow can then deploy in many places (Docker, Azure ML, Spark UDF)
- Open interface allows easy integration from the community

**Key enabler: built around  
Programmatic APIs, REST APIs & CLI**

## Modular design

- Allow different components individually (e.g., use MLflow’s project format but not its deployment tools)
- Not monolithic
- But Distinctive and Selective

**Key enabler: distinct components  
(Tracking/Projects/Models/Registry)**

# MLflow Components

## mlflow Tracking

Record and query experiments: code, data, config, and results

## mlflow Projects

Package data science code in a format that enables reproducible runs on any platform

## mlflow Models

Deploy machine learning models in diverse serving environments

new

## mlflow Model Registry

Store, annotate and manage models in a central repository

[databricks.com/  
mlflow](https://databricks.com/mlflow)



[mlflow.org](https://mlflow.org)



[github.com/mlflow](https://github.com/mlflow)



[twitter.com/MLflow](https://twitter.com/MLflow)

# Key Concepts in MLflow Tracking

**Parameters:** key-value inputs to your code

**Metrics:** numeric values (can update over time)

**Tags and Notes:** information about a run

**Artifacts:** files, data, and models

**Source:** what code ran?

**Version:** what of the code?

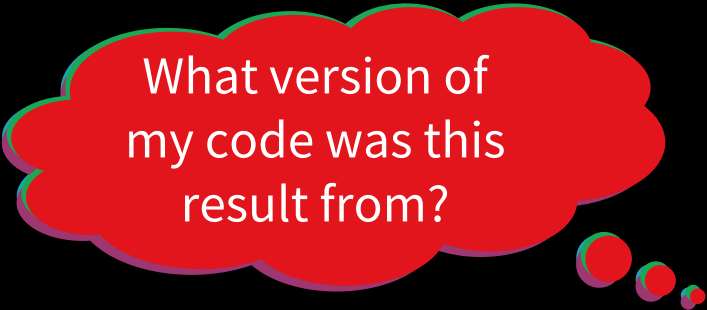
# Model Development without MLflow

```
data    = load_text(file)
ngrams  = extract_ngrams(data, N=n)
model   = train_model(ngrams,
                      learning_rate=lr)
score   = compute_accuracy(model)

print("For n=%d, lr=%f: accuracy=%f"
      % (n, lr, score))

pickle.dump(model, open("model.pkl"))
```

```
For n=2, lr=0.1: accuracy=0.71
For n=2, lr=0.2: accuracy=0.79
For n=2, lr=0.5: accuracy=0.83
For n=2, lr=0.9: accuracy=0.79
For n=3, lr=0.1: accuracy=0.83
For n=3, lr=0.2: accuracy=0.82
For n=4, lr=0.5: accuracy=0.75
...
```



What version of  
my code was this  
result from?

# MLflow Tracking API: *Simple & Pythonic!*

**mlflow**  
Tracking

Record and query  
experiments: code,  
configs, results,  
...etc

```
import mlflow
import mlflow.tensorflow

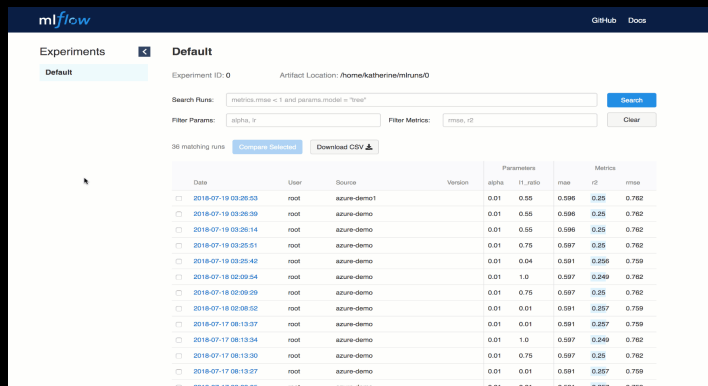
# log model's tuning parameters
with mlflow.start_run() as run:
    mlflow.log_param("layers", layers)
    mlflow.log_param("alpha", alpha)

# log metrics and model
mlflow.log_metric("mse", model.mse())
mlflow.log_artifact("plot", model.plot(test_df))
mlflow.tensorflow.log_model(model)
```

# Model Development *with* MLflow is *Simple*!

```
data    = load_text(file)
ngrams  = extract_ngrams(data, N=n)
model   = train_model(ngrams,
                      learning_rate=lr)
score   = compute_accuracy(model)
with mlflow.start_run() as run:
    mlflow.log_param("data_file", file)
    mlflow.log_param("n", n)
    mlflow.log_param("learn_rate", lr)
    mlflow.log_metric("score", score)
    mlflow.sklearn.log_model(model)
```

```
$ mlflow ui
```



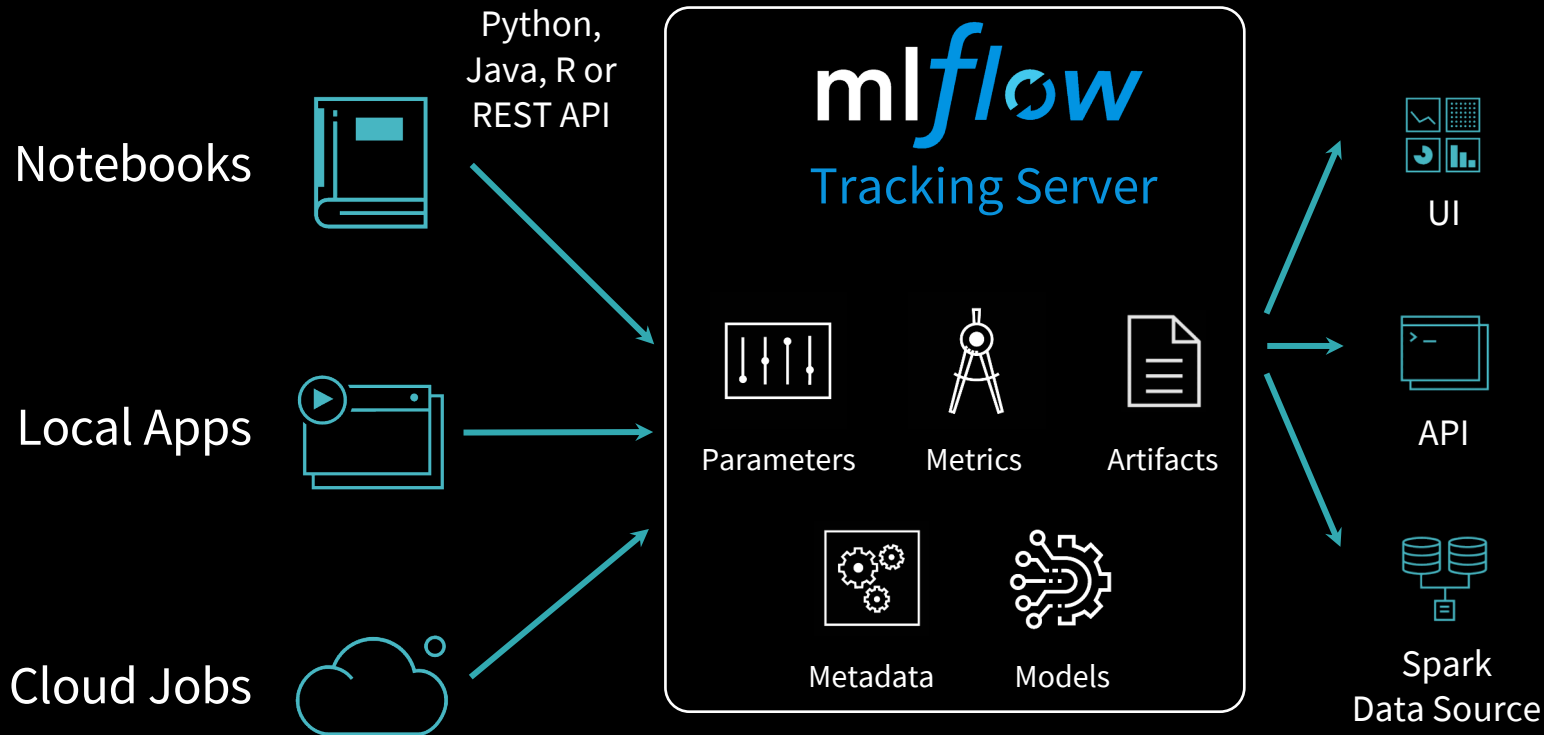
MLflow UI Screenshot: Experiments page showing a table of runs.

Date	User	Source	Version	params.alpha	params.l1_ratio	metrics.max	metrics.f1	metrics.mse
2018-07-19 03:05:53	root	azure-demo1	0.01	0.55	0.595	0.25	0.762	
2018-07-19 03:06:39	root	azure-demo	0.01	0.55	0.595	0.25	0.762	
2018-07-19 03:06:14	root	azure-demo	0.01	0.55	0.595	0.25	0.762	
2018-07-19 03:05:51	root	azure-demo	0.01	0.75	0.587	0.25	0.762	
2018-07-19 03:05:42	root	azure-demo	0.01	0.04	0.591	0.255	0.759	
2018-07-19 02:09:54	root	azure-demo	0.01	1.0	0.597	0.249	0.762	
2018-07-18 02:09:29	root	azure-demo	0.01	0.75	0.597	0.25	0.762	
2018-07-18 02:08:52	root	azure-demo	0.01	0.01	0.591	0.257	0.759	
2018-07-17 08:13:27	root	azure-demo	0.01	0.01	0.591	0.257	0.759	
2018-07-17 08:13:34	root	azure-demo	0.01	1.0	0.597	0.249	0.762	
2018-07-17 08:13:30	root	azure-demo	0.01	0.75	0.597	0.25	0.762	
2018-07-17 08:13:27	root	azure-demo	0.01	0.01	0.591	0.257	0.759	
2018-07-17 08:08:05	root	azure-demo	0.01	0.01	0.591	0.257	0.759	

Track parameters, metrics,  
output files & code version

Search using UI or API

# MLflow Tracking



```
$ export MLFLOW_TRACKING_URI <URI>  
mlflow.set_tracking_uri(URI)
```



# MLflow Tracking Backend Stores

## 1. Entity (Metadata) Store

- FileStore (local filesystem)
- SQLStore (via SQLAlchemy)
  - PostgreSQL, MySQL, SQLite

## 2. Artifact Store

- S3 backed store
- Azure Blob storage
- Google Cloud storage
- DBFS artifact repo


# What Did We Talk About?

## **Modular Components greatly simplify the ML lifecycle**

- Open machine learning platform
- Available APIs: Python, Java & R (Soon Scala)
- *Simple. Modular. Easy-to-use.*
- *Offers positive developer experience to get started!*

# Project Contributors over Time

Today at 191 contributors



# Learning More About MLflow

- `pip install mlflow` to get started
- Find docs & examples at [mlflow.org](https://mlflow.org)
- <https://github.com/mlflow/mlflow>
- [tinyurl.com/mlflow-slac](https://tinyurl.com/mlflow-slac)
- [dbricks.co/mlflow-tutorials](https://dbricks.co/mlflow-tutorials)

# MLflow Tracking Tutorials

<https://github.com/dmatrix/mlflow-workshop-part-1>



# Thank you! ☺

Q & A

[jules@databricks.com](mailto:jules@databricks.com)

[@2twitme](#)

<https://www.linkedin.com/in/dmatrix/>